

# Verification of Solana Programs

Jorge A. Navas (Certora) and  
Arie Gurfinkel (University of Waterloo and Certora)


Venice, May 26th 2023

Symposium on Challenges of Software Verification (CSV)



# DeFi in one slide

- Economic process completely defined by code
- Fairly complex code
- Examples
  - Lending
  - Exchange
  - Options
  - Auctions
- 50 Billion dollars in the bear market



## Interesting DeFi Bugs 2022/3

- **Euler Finance \$200M** – DonateToReserves() function didn't check for account debt health, allowing for bad debt to accrue and for the collateral to be liquidated at a large discount to the attacker
- **Yearn Finance V1 \$10M** – Misconfiguration of one of the underlying asset addresses in the USDT pool allowed an attacker to drain the whole vault
- **Safemoon \$9M** – Upgraded contract didn't use access control for the burn() function. The attacker burned tokens from the Safemoon pool on a DEX, inflated the price and sold tokens into the pool
- **Platypus \$8.5M** – EmergencyWithdraw() didn't check for debt, so the attacker could take max loan for his collateral, and then simply emergency withdraw the collateral
- **Hundred \$7.4M** – "First depositor" bug where the attacker could manipulate the exchange rate and borrow way more than allowed

# Why Formally Verify DeFi?



Code is law



Billions of dollars at stake



Code is typically medium-size/modular



But bugs are hard to find  
**Happens in rare scenarios**



New code is produced frequently



**Maker**  
@MakerDAO

UPDATE ON MULTI-COLLATERAL DAI:  
The code is ready and formally verified. The first time ever  
a major dapp has been formally verified.  
Learn more: [medium.com/makerdao/the-c...](https://medium.com/makerdao/the-c...)  
[#FormalVerification](#) [#DAI](#) [\\$DAI](#) [\\$MKR](#) [#MKR](#)

12:07 AM · Sep 18, 2018



**Lido**  
@LidoFinance

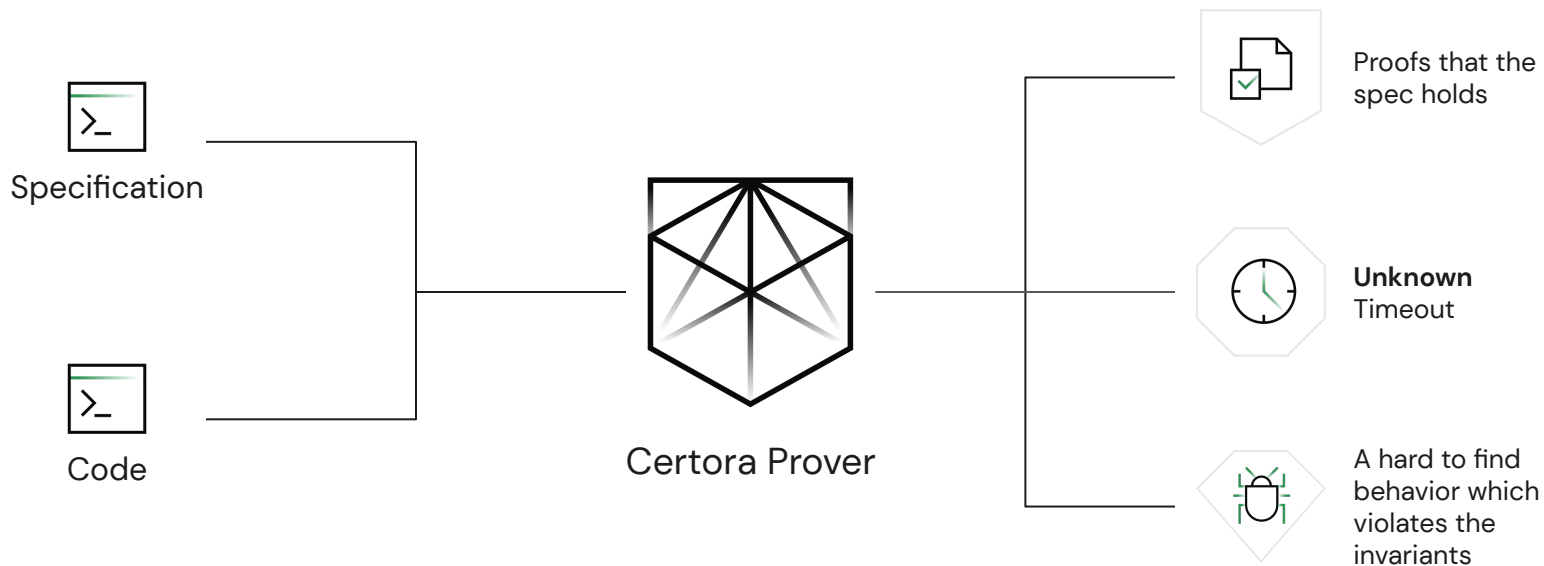
The Lido-on-Ethereum protocol team is doing all it can to  
make sure the protocol upgrade is secure and issue-free,  
including conducting thorough security audits, performing  
formal verification, and extensively testing on Goerli.

9:01 PM · Feb 28, 2023 · **1,951** Views



certora

# The Certora Approach: Automatic Formal Verification







# Critical Bugs Found by Certora Prover

## Solvency

- If everybody runs to the bank  
Bank still fulfills all commitments
- Users' money cannot  
be locked or lost

## Bugs prevented by the Certora-Prover **missed in manual audits** by top auditors









 SushiSwap	\$807M	 AAVE	\$6.5B	 Compound	\$2.7B	 Balancer	\$1.18B
Strategy	2	V3	1	Comet	5	V2	2
Trident	5	V2	2	V2	5		
KashiPair	3						
DutchAuction	1						



Solidity  
[@solidity\\_lang](#)

"We thank all contributors who made this release possible.  
Special thanks goes to [@johnadtoman](#) of [@Certoralnc](#) for  
reporting the inline assembly memory side effects bug!"

# Why Formally Verify Solana (<https://solana.com>) ?

#	Name	Protocols	1d change(TVL)	1w change(TVL)	1m change(TVL)	TVL	Mcap	Mcap/TVL
1	 Ethereum ETH	810	▲ 0.38%	▼ 3.76%	▼ 12.83%	\$51.48B	\$220.16B	4.28
2	 Tron TRON	32	▲ 0.05%	▼ 2.99%	▼ 2.89%	\$5.40B	\$6.26B	1.16
3	 BSC BNB	612	▲ 0.19%	▼ 2.04%	▼ 7.26%	\$5.21B	\$48.51B	9.31
4	 Arbitrum ARB	337	▲ 0.18%	▼ 1.76%	▲ 2.59%	\$2.68B	\$1.48B	0.55
5	 Polygon MATIC	426	▲ 0.89%	▼ 2.55%	▼ 12.48%	\$1.21B	\$7.86B	6.49
6	 Optimism OP	144	▲ 0.94%	▼ 4.01%	▼ 11.16%	\$1.00B	\$540.39M	0.54
7	 Avalanche AVAX	318	▲ 0.47%	▼ 2.62%	▼ 15.41%	\$979.23M	\$4.95B	5.05
8	 Solana SOL	114	▼ 0.53%	▼ 2.55%	▼ 16.62%	\$514.09M	\$8.17B	15.9

EVM

Non-EVM

# Why Formally Verify Solana?

- Benefits:
  - Based on general purpose programming languages: Rust, C/C++
  - Reusing existing eBPF virtual machine:
    - Support multiple (or even combination of) input languages
  - Programs are stateless: all data is passed as function arguments
    - Non-interference (easier to shard)
- Challenges:
  - Verification of low-level eBPF/SBF is harder
  - No common format between apps (data format is up to the app):
    - Inputs are just array of bytes
    - Serialization/deserialization
  - Compiled Rust can be harder to verify than human-written C
    - Rust union types, dangling pointers, etc.



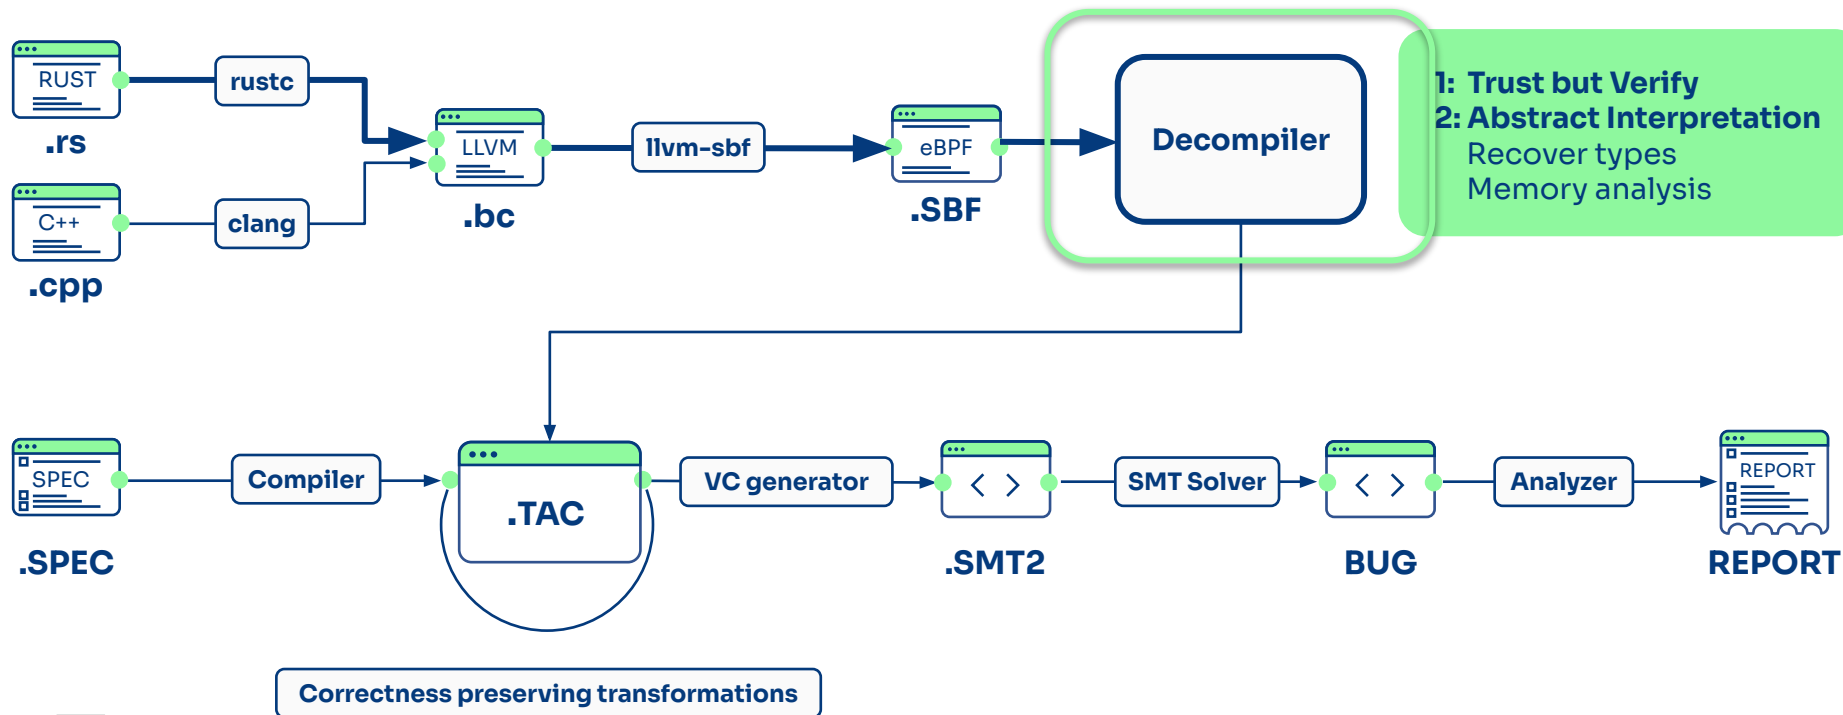
# Solana Programming (not in this talk)

- **Accounts**
  - Fields: lamports, owner, executable, data, rent epoch
  - Program and Data accounts
- **Transactions** consist of **instructions**
- All programs are *stateless*: any data they interact with is stored in separate accounts that are passed in via instructions
- **PDA**s (Program Derived Address): data account owned by programs instead of users
  - Used to implement associative maps
- **CPI** (Cross Program Invocations)
- **Deserialize/Serialize**

<https://solanacookbook.com/>

# Certora Prover Architecture for Solana

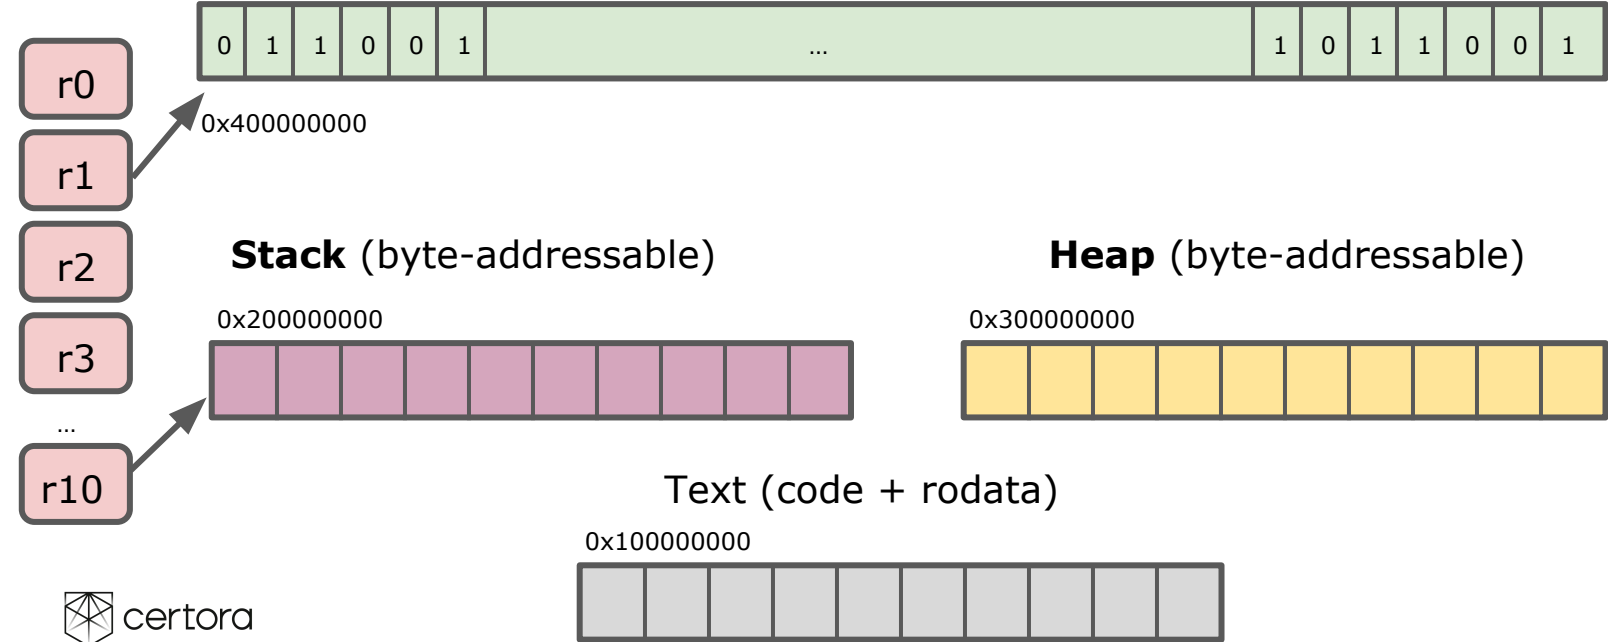
In this talk



# eBPF/SBF Virtual Machine

## Blockchain State (program inputs)

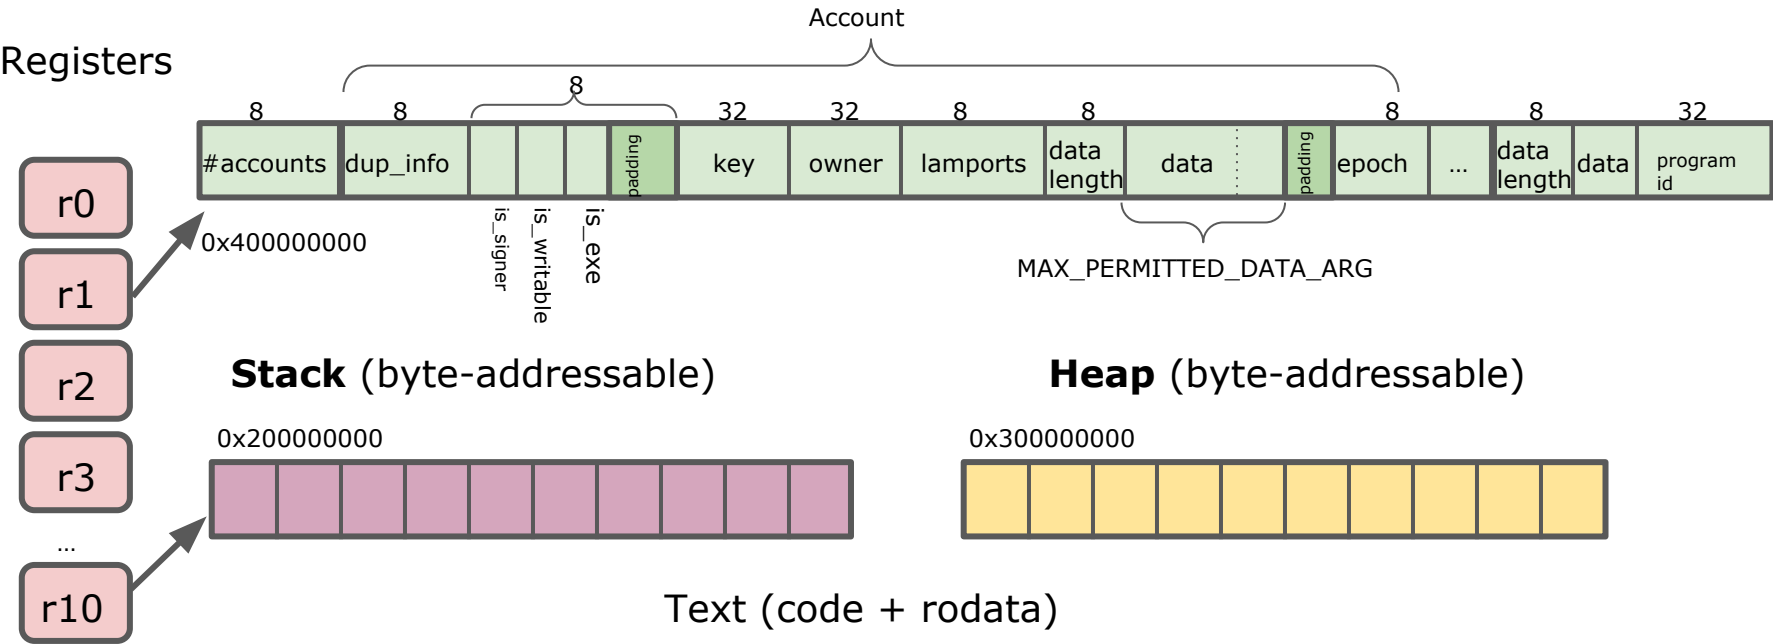
Registers



# eBPF/SBF Virtual Machine

## (Deserialized) Blockchain State

Registers



# SBF Instruction Set

- Currently, three different dialects with similar bytecodes: bpf/sbf/sbfv2
- RISC-like instruction set
- 11 general-purpose, 64-bit registers
  - r10 is read-only frame pointer to access to stack
- ALU, JUMP, LOAD, STORE, MOVE
  - Jumps use only relative constant offsets: CFG construction is decidable
- Syscalls and eBPF-to-eBPF (internal) calls
  - r0: return
  - r1, ..., r5: caller-saved (volatile) registers
  - r6, ..., r9: callee-saved (non-volatile) registers
- No type information: no distinction between numbers and pointers
- Direct and indirect function calls: call graph construction is undecidable

# SBF Disassembler


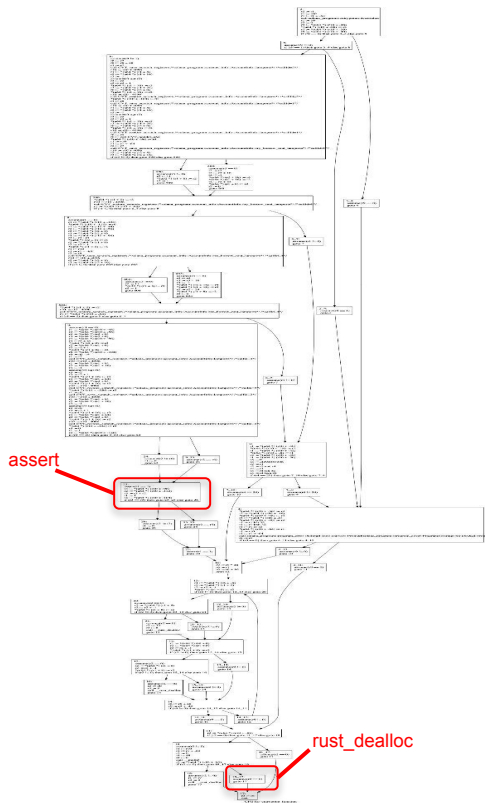
1. Translate ELF to a sequence of three-address instructions
  - Resolve Solana-specific relocations
2. CFG and Call graph construction: one per function
  - Indirect calls not supported
3. Inline all internal functions
  - Explicit modeling of call semantics
4. Compute Cone-of-Influence and slice program
5. Memory analysis
6. Translation to TAC program

# Memory Analysis Assumptions

The analysis is sound under the following assumptions:

1. Memory safety
  - Absence of out-of-bounds accesses
  - Stack/Heap/Blockchain memory is initialized
2. First read from blockchain state returns non-deterministic values
  - Pointers do not alias with any other pointer
3. Each memory read accesses the same number of bytes last written
  - Checked by the analysis

# Rust compiles to large programs



- Many irrelevant paths:
  - error paths
  - free pointers
- We only care about paths that can influence the evaluation of assertions

Solution: dataflow analysis that removes any path that is not in the Cone-Of-Influence (CoI)



# Rust enum types

```
pub fn process_withdraw(  
    program_id: &Pubkey,  
    accounts: &[AccountInfo],  
    amount: u64,  
    expected_decimals: u8,  
    new_decryptable_available_balance: DecryptableBalance,  
    proof_instruction_offset: i64,  
) -> ProgramResult {
```

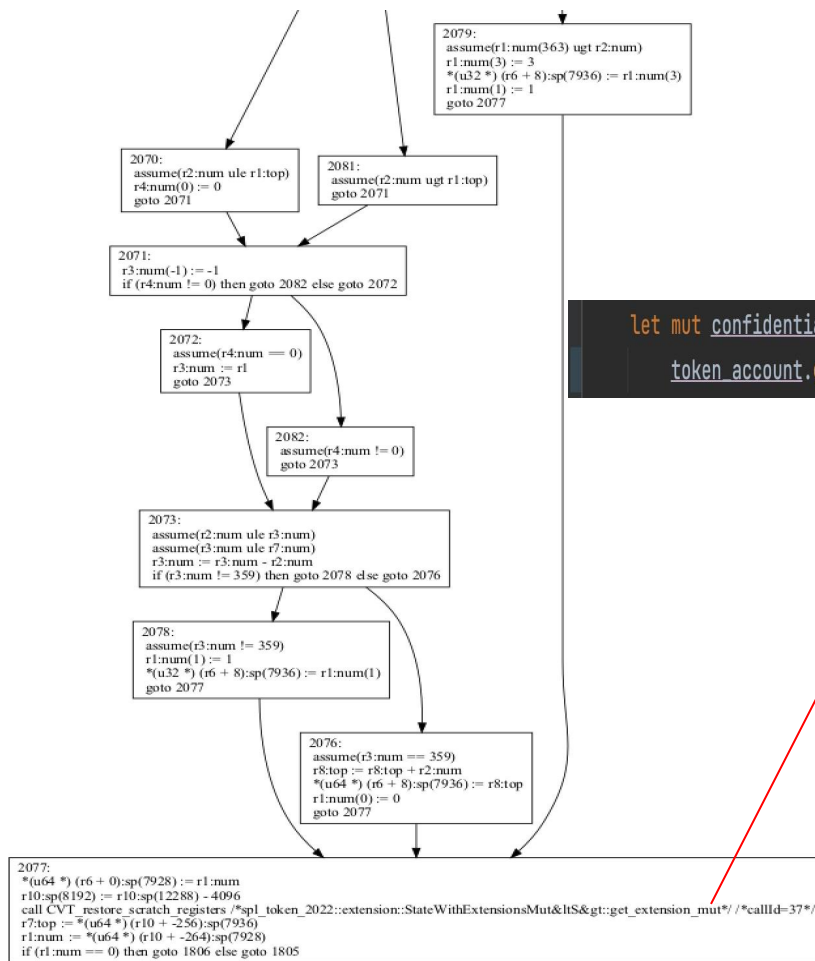
Result<(), ProgramError>

Result<&mut V, ProgramError>

```
let mut confidential_transfer_account : &mut ConfidentialTransferAccount =  
    token_account.get_extension_mut:<ConfidentialTransferAccount>()?;
```

question mark (?) operator

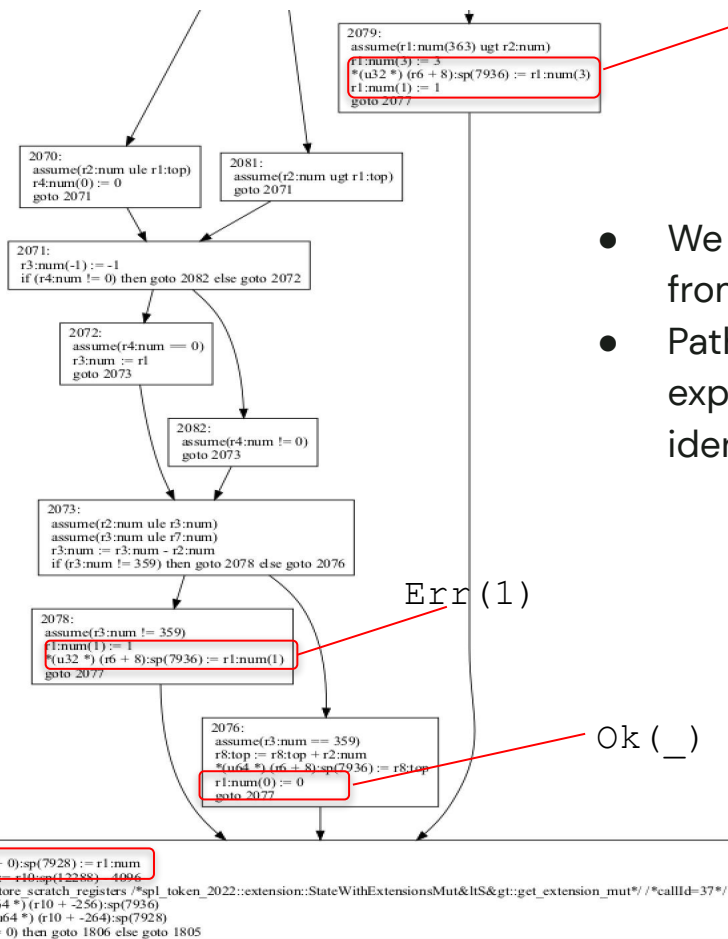
# Rust enum types



```
let mut confidential_transfer_account : &mut ConfidentialTransferAccount =
    token_account.get_extension_mut::<ConfidentialTransferAccount>()?;
```

# Rust enum types

*r1* is the  
discriminant  
0 = Ok  
1 = Err



Err (3)

- We need to discriminate *error* from *ok* paths
- Path-sensitive analysis is expensive and it is not easy to identify the discriminants

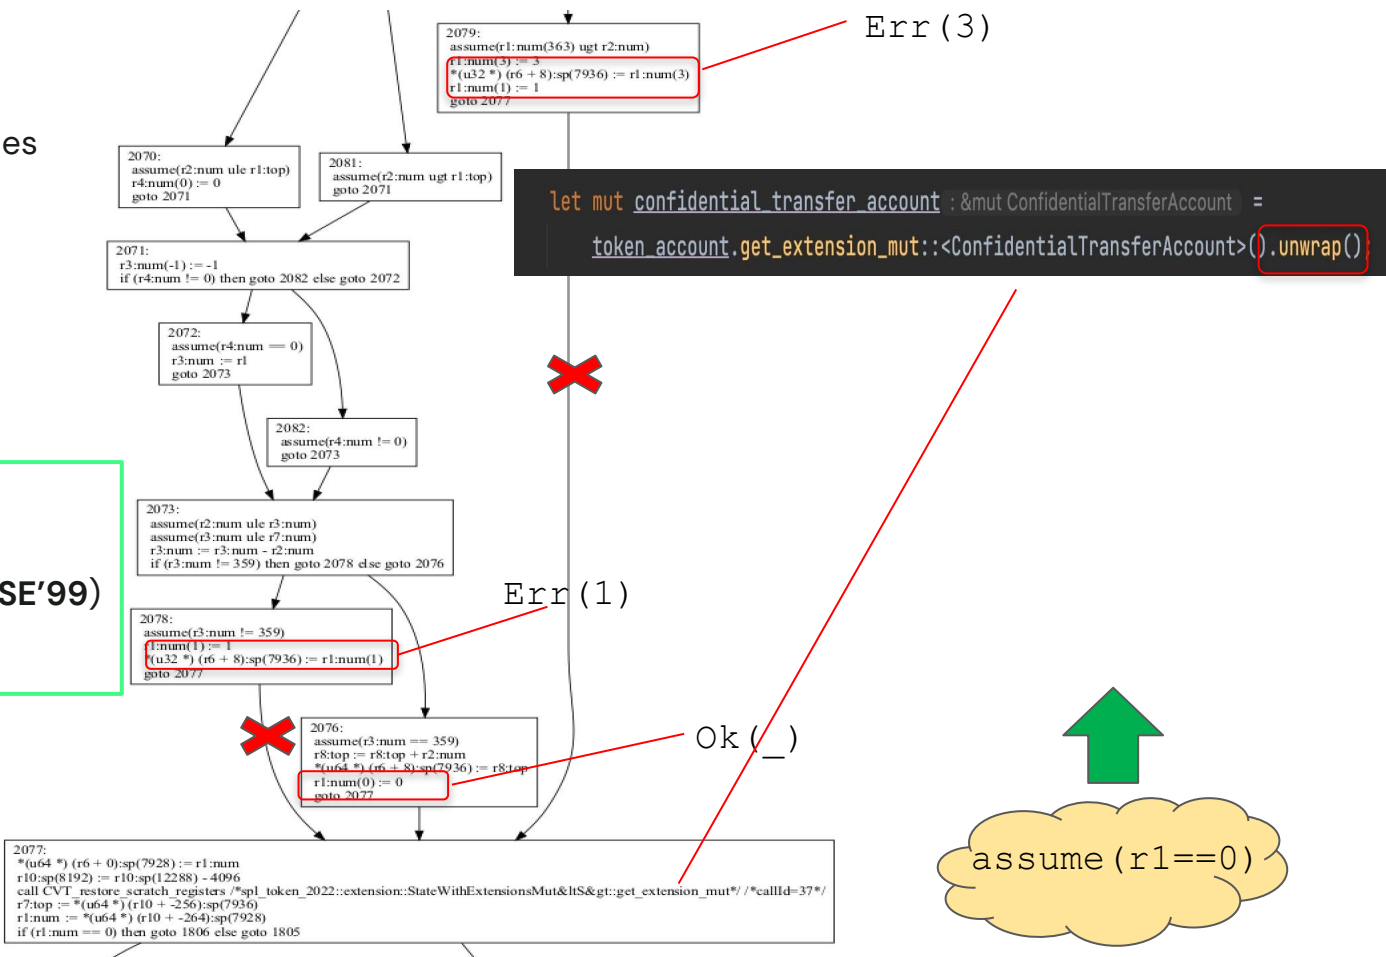
Err (1)

Ok ( \_ )

# Rust enum types

We typically prove properties under the assumption that functions return `Ok`

Solution: iterative forward+backward analysis (Cousot&Cousot JLP'92/ASE'99) to prune error paths

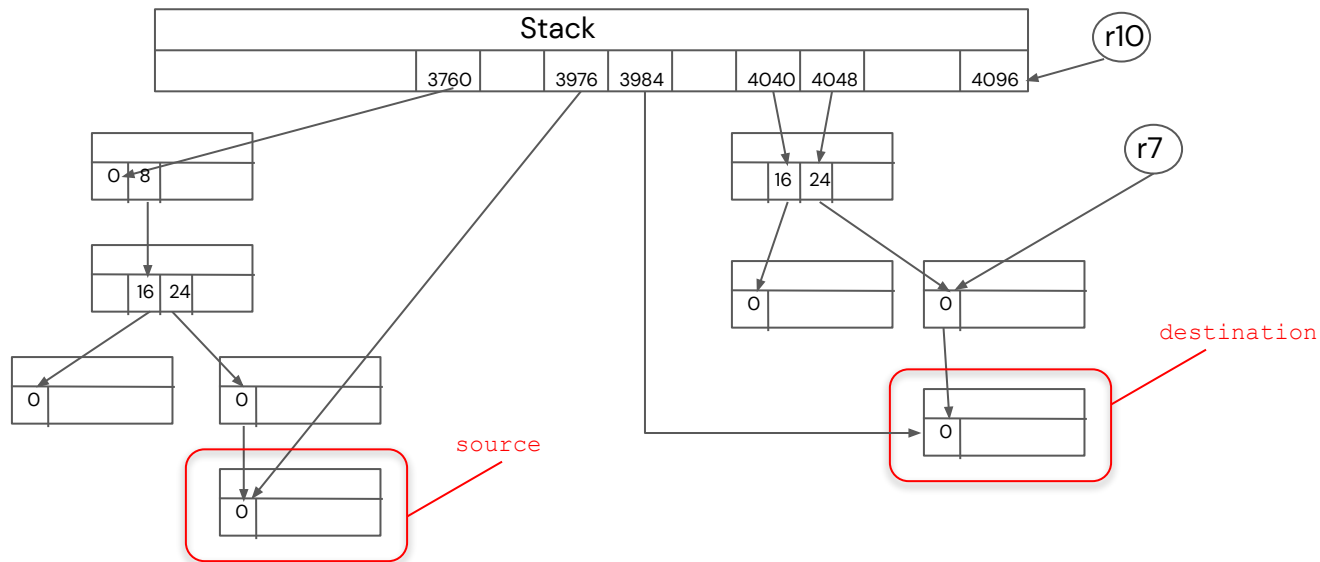


# Analysis of SBF code

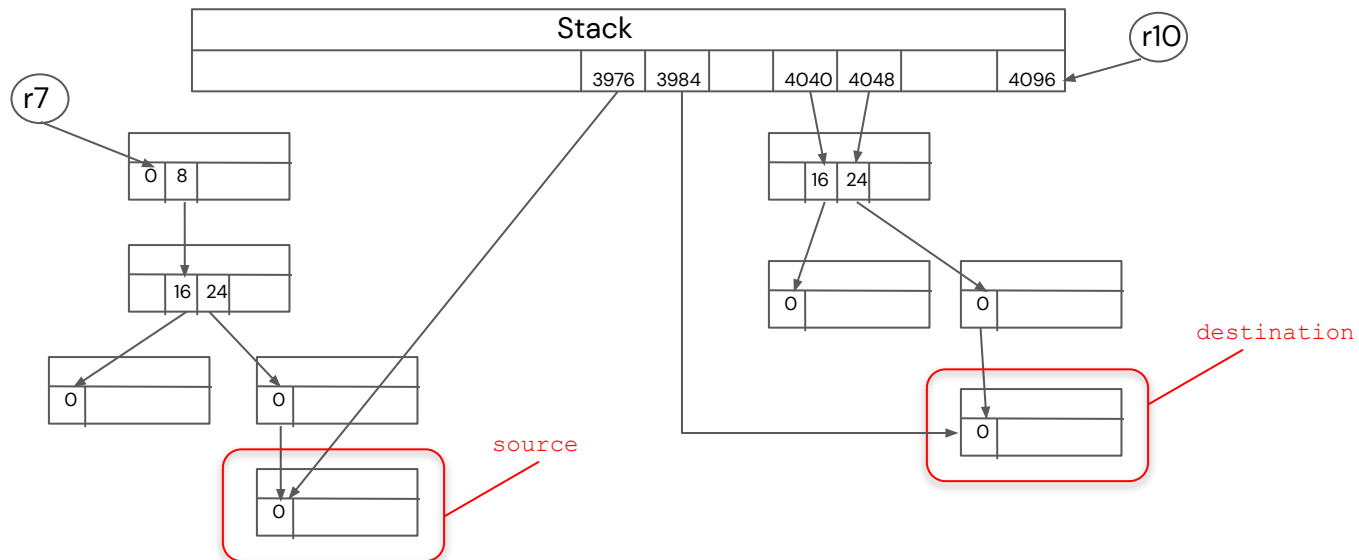
- Disassembler needs to translate SBF into a TAC program **without side effects**
  - TAC memory operations have an explicit argument “mem” that represents the (possibly infinite) set of memory locations being accessed
  - Two TAC memory ops do not alias if they have different “mem” names
- How: static memory partitioning
  - Split all program memory (stack, heap, and inputs) into a **finite** set of **disjoint** regions
  - For each memory instruction, map the memory location to a region
- Challenges:
  - No explicit allocation sites for program inputs because they are allocated either before the SBF program is loaded or by deserialization
  - Strong vs weak updates

# Analysis of SBF code

- Solution 1: flow-insensitive/field-sensitive pointer analysis (**Gurfinkel&Navas SAS'17**)
  - Adopted in LLVM-based verifiers such as SeaHorn and SMACK
  - Easy to model in SMT: one single points-to graph for the whole program

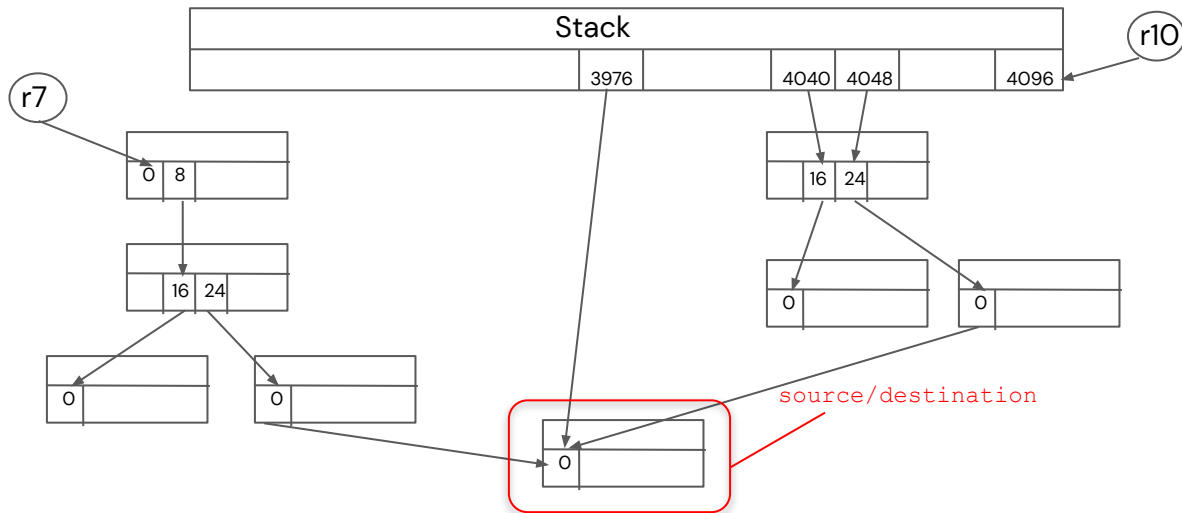


# Analysis of SBF code



- Registers must be tracked flow-sensitively
  - They can be re-assigned at each instruction

# Analysis of SBF code



- Registers must be tracked flow-sensitively
  - They can be re-assigned at each instruction

- Stack must be tracked flow-sensitively

- LLVM back-end reuses stack allocations

```
lifetime.start.p0i8(%p1)
call try_borrow_mut_lamports(%p1) // %p1 points to src
lifetime.end.p0i8(%p1)
lifetime.start.p0i8(%p2)
call try_borrow_mut_lamports(%p2) // %p2 points to dst
lifetime.end.p0i8(%p2)
```

- Same slot 3976 in SBF for %p1 and %p2



# Analysis of SBF code

- Solution 2: flow-sensitive pointer analysis
  - Solution adopted by verifiers such as Predator
  - Very precise but expensive: one points-to graph per basic block
  - Harder to model in SMT: a memory instruction can use different “mem” depending on which predecessor reaches the instruction

# Analysis of SBF code

- Our solution:
  - Flow-sensitive stack and registers
  - Flow-insensitive heap and program inputs
  - **Stack scalarization:**
    - Each stack slot is translated to a scalar variable
    - This allows **strong updates** on local variables
    - Precise and easy to model in SMT
  - Weak updates on heap and program inputs
    - Still easy to model in SMT

# Conclusions

- Solidity/EVM has attracted most of the attention of the verification community
- Verification of Solana contracts is a very exciting new research area
- Based on thrilling Rust and eBPF technology
  - A lot of the ideas and solutions can be reused in different contexts
- Both (compiled) Rust and SBF pose unique challenges to verification
- **Certora is building the first automatic verifier for Solana contracts!**

## Many challenges are still to solve ...

- Solana
  - a. Cross-program invocations (CPI)
  - b. Automatic handling of serialization/deserialization
  - c. Verifying multiple transactions/instructions
    - For now, we focus on one instruction at the time, and manually provide context invariants
    - However, most exploited vulnerabilities used multiple instructions and transactions
  - d. Fuller model of transaction state
    - e.g., support instruction introspection (heavily used for implementing confidentiality)
  - e. Richer model of the blockchain environment: e.g., PDA-based links between accounts
- Rust/SBF
  - a. More precise memory abstraction to support Rust enum types
  - b. More precise abstractions for the heap (e.g., Box, Vec, ...)
- SMT
  - a. Improve domain-specific treatment of non-linear arithmetic

# Demo: SPL Token 2022

<https://spl.solana.com/token-2022>

# Verification harness for `process_withdraw`

```
64 fn cvt_harness_process_withdraw(  
65     program_id: &Pubkey,  
66     _accounts: &[AccountInfo],  
67     _instruction_data: &[u8],  
68 ) -> ProgramResult {  
69     // Create non-deterministic inputs for process_withdraw  
70     let token_account_info = CVT_nondet_account_info();  
71     let mint_account_info = CVT_nondet_account_info();  
72     let instructions_sysvar_info = CVT_nondet_account_info();  
73     let authority_info = CVT_nondet_account_info();  
74     let acc_infos : [?; 4] = [token_account_info.clone(), mint_account_info.clone(), instructions_sysvar_info.clone(), authority_info.clone()];  
75     let amount = CVT_nondet_u64();  
76     let expected_decimals = CVT_nondet_u8();  
77     let new_decryptable_available_balance: DecryptableBalance = cvt_confidential::CVT_mk_decryptable_balance();  
78     let proof_instruction_offset = CVT_nondet_i64();  
79     // Call the function under verification  
80     process_withdraw(program_id, accounts: &acc_infos, amount, expected_decimals, new_decryptable_available_balance, proof_instruction_offset).unwrap();  
81     // Check the property  
82     let token_account_data : &? = &acc_infos[0].data.borrow();  
83     let token_account : StateWithExtensions<Account> = StateWithExtensions::<Account>::unpack( input: token_account_data).unwrap();  
84     let confidential_transfer_account : &ConfidentialTransferAccount =  
85         token_account.get_extension::<ConfidentialTransferAccount>().unwrap();  
86     cvt::CVT_assert(confidential_transfer_account.encryption_pubkey == cvt_confidential::get_proof_withdraw_account().pubkey);  
87     Ok()  
88 }
```

Non-deterministic inputs

function under verification

property

# process\_withdraw

```
443 // ZK-proof certifies that the account has enough available balance to withdraw the amount.
444 let zkp_instruction : Instruction =
445     get_instruction_relative( index_relative_to_current: proof_instruction_offset,
446                               instruction_sysvar_account_info: instructions_sysvar_info)?;
447 let proof_data : &WithdrawData = decode_proof_instruction:::<WithdrawData>(
448     expected: ProofInstruction::VerifyWithdraw,
449     &zkp_instruction,
450 )?;
451
452 // Check that the encryption public key associated with the confidential extension is
453 // consistent with the public key that was actually used to generate the zkp.
454 // ===== THIS CHECK WAS ADDED BY AUDITOR =====
455 // SEE https://github.com/solana-labs/solana-program-library/pull/3768/
456 if confidential_transfer_account.encryption_pubkey != proof_data.pubkey {
457     return Err(TokenError::ConfidentialTransferElGamalPubkeyMismatch.into());
458 }
459
460 // Prevent unnecessary ciphertext arithmetic syscalls if the withdraw amount is zero
461 if amount > 0 {
462     confidential_transfer_account.available_balance =
463         syscall::subtract_from( ciphertext: &confidential_transfer_account.available_balance, amount)
464         .ok_or( err: ProgramError::InvalidInstructionData).unwrap();
465 }
466
467 // Check that the final available balance ciphertext is consistent with the actual ciphertext
468 // for which the zero-knowledge proof was generated for.
469 if confidential_transfer_account.available_balance != proof_data.final_ciphertext {
470     return Err(TokenError::ConfidentialTransferBalanceMismatch.into());
471 }
```

ZK-Proof

Bug source

# Mocking process\_withdraw

```
443 // ZK proof certifies that the account has enough available balance to withdraw the amount
444 let proof_data : &WithdrawData = cvt_confidential::get_proof_withdraw_account();
445
446 // Check that the encryption public key associated with the confidential extension is
447 // consistent with the public key that was actually used to generate the zkp.
448 // ===== THIS CHECK WAS ADDED BY AUDITOR =====
449 // SEE https://github.com/solana-labs/solana-program-library/pull/3768/
450 if confidential_transfer_account.encryption_pubkey != proof_data.pubkey {
451     return Err(TokenError::ConfidentialTransferElGamalPubkeyMismatch.into());
452 }
453
454 // Prevent unnecessary ciphertext arithmetic syscalls if the withdraw amount is zero
455 if amount > 0 {
456     confidential_transfer_account.available_balance =
457         syscall::subtract_from( ciphertext: &confidential_transfer_account.available_balance, amount)
458         .ok_or( err: ProgramError::InvalidInstructionData).unwrap();
459 }
460
461 // Check that the final available balance ciphertext is consistent with the actual ciphertext
462 // for which the zero-knowledge proof was generated for.
463 if confidential_transfer_account.available_balance != proof_data.final_ciphertext {
464     return Err(TokenError::ConfidentialTransferBalanceMismatch.into());
465 }
466
467 confidential_transfer_account.decryptable_available_balance = new_decryptable_available_balance;
```

mock



# Usage of Certora Prover

- Compile Solana Program to generate SBF

```
% cargo build-sbf --arch=sbfv2
Finished release [optimized] target(s) in 6.24s
```

- Run Certora Prover on buggy version:

```
% certoraRun.py target/sbf-solana-solana/release/spl_token_2022.so
Violated: cvt_harness_process_withdraw
cvt_harness_process_withdraw: A property is violated
Done Sat May 20 14:14:45 MDT 2023
Reports in file:///Users/jorge/SPL/token/program-2022/src/certora_files/emv-1-
release-spl_token_2022.so-20-May--14-14/Reports
```

- Run Certora Prover on fixed version

(<https://github.com/solana-labs/solana-program-library/pull/3768/>):

```
% certoraRun.py target/sbf-solana-solana/release/spl_token_2022.so
Verified: cvt_harness_process_withdraw
cvt_harness_process_withdraw: Properties successfully verified on all inputs
```

